




So what does a *realistic* environment look like?



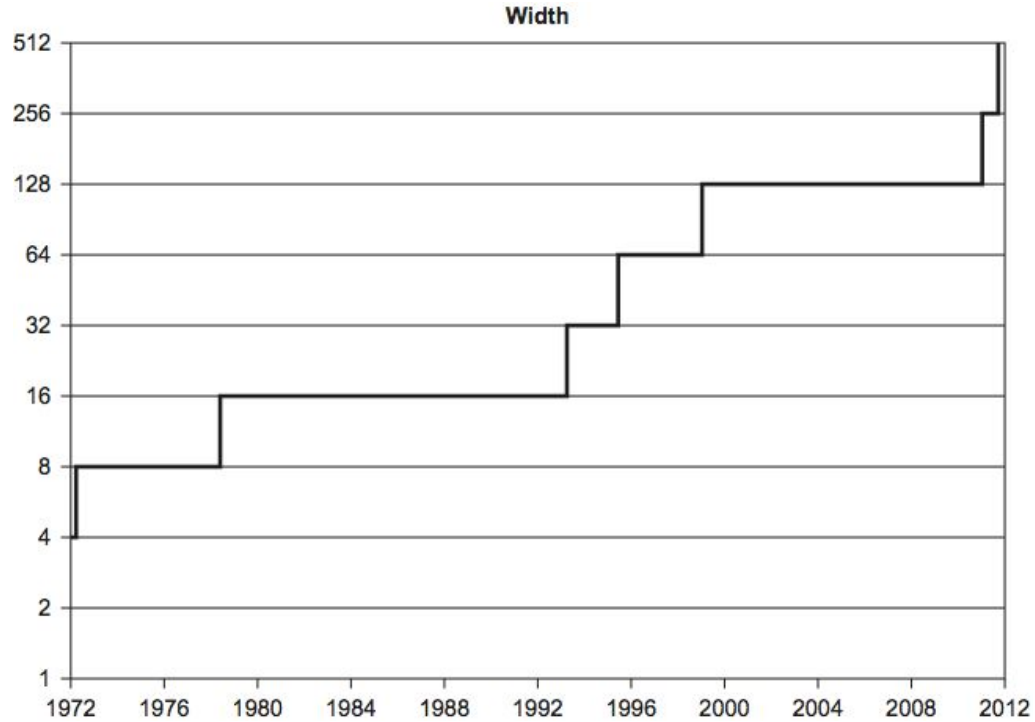
# Moore's Law



# More Cores



# Wider Vector Units




# Consequences for Developers

- free lunch is over / disruptive
  - software used to speed up with clock
  - to make free lunch work again, software now needs to scale
- Efficient Cache Utilization  
restructure software to more predictable access patterns
- Vectorization  
use vector units for SIMD work

# This Course: more Parallelism

- 2 CPUs
- 8 Cores, 16 Threads
- 256 GB RAM (NUMA)
- 4 Xeon Phi Cards
- 60 cores, 240 Threads
- 8GB RAM
- $(4.04 + 0.2)$  TFlop/s



# Course outlook

- Programming of highly parallel machines
  - OpenMP
  - MPI
  - threading libraries (TBB)
  - (OpenCL)
- Profiling and Debugging
- Performance Optimizations
- Applications


# History

- 2008 Intel announced GPGPU
- Larrabee - was never produced
- 2010 MIC architecture announced  
*Many Integrated Core*
- 2012 First Cards shipped under the name  
Xeon Phi (Knights Corner)
- 2015 Next MIC Generation  
Xeon Phi (Knights Landing) 3 TFlop

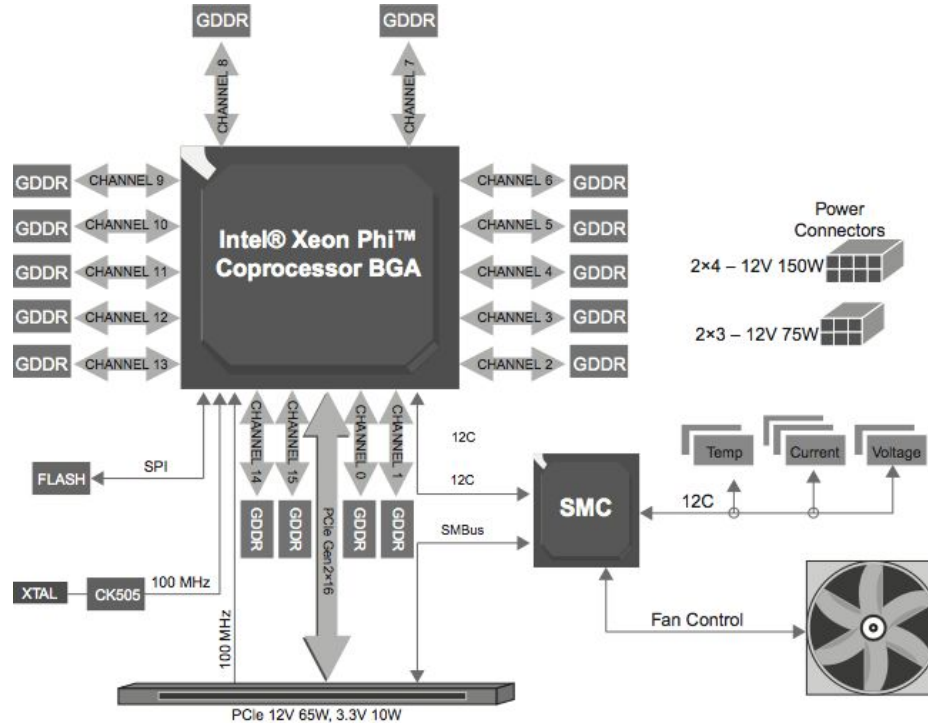


# Xeon Phi Architecture


- accelerator card connected via the PCIe bus
- „motherboard“ for a computer with 60 cores and 8 GB RAM
- flash memory with boot routine that boots a small linux system from the host System
- emulates a network interface over PCIe



# Xeon Phi Architecture




# Chip Layout



# Core Structure

- 64-bit architecture
- in-order execution
- favors multithreading over complex core
- instruction decode has latency of 2 cycles! →  
Single threaded (compute-bound) code will only reach 50% peak! 512-bit VPU



# Hardware Threads

- Every core has replicated state for 4 hardware Threads
- Similar to Hyperthreading(HT) on regular x86 Cores
- Difference: while HT might improve performance by about 5-10% (sometimes performance even drops), Hardware-Threads on the Xeon Phi are **always** beneficial! (use at least 2!)
- scheduled in **round-robin** fashion

# Cache organization

- 31kb L1-data and 31kb L1-Instruction Cache
- 512kb L2 (last-level) cache
- **fully coherent** among the cores
- total cache of Xeon Phi Die: 30 MB
  - software may use everything between 512kb and 30MB as cache memory
- L1 latency: 1 cycle, L2 latency: 11 cycles
- Translation-Lookaside-Buffer (TLB) supports **huge pages** of 2MB

# Prefetching

- a request to the Cache/Memory-Access Subsystem
- look ahead and bring the data that is expected to be used into L1/L2 cache for the near future
- fetching memory is slow → **hide latency!**
- hardware prefetching & software prefetching
- software prefetching either manually or automatically by the compiler
- can be scheduled simultaneously to VPU instruction

# Vector Processing Unit

- Each core has a VPU of 512-bit width
- similar to SSE/AVX but not backwards compatible
- can process 16 single-/8 double-precision floats in a single instruction
- at least as important to scaling as scaling to Threads (SIMD workload)
- Special support for transcendental functions, **fused-multiply-add** (FMA)  
 $c += a * b$



# Performance Metrics

- Floating Point Operations / second  
Throughput: 1 **retired** operation / cycle  
1.053 GHz: 1053000000 instr./second  
8-unit wide vector instructions  
2 Flops per unit (FMA)  
60 cores  
→  $1.053 * 10^9 * 8 * 2 * 60 = 1011 \text{ GFlop/s}$
- Memory Bandwidth  
8 Memory Controllers  
GDDR5 -> dual Channel, 4 Bytes per Transfer  
5.0 GT/s ( $5 * 10^9$  transfers per second)  
-> 320 GB/s (**realistic ~60% of this peak**)

# Reality of Parallelization

- in an ideal world **n-way** multiprocessor provides an n-fold of computational power
- in practice this almost never happens
- often parallelization comes with **communication and coordination** costs

# Example

- Five friends want to paint a five-room house
- all rooms are the same size and everyone paints at about the same rate  
→ low coordination/communication: 5x speedup
- What if the rooms are of different size?  
Same example, but: one room is twice as big as the others  
→ dominates the overall time
- adding more painters incurs coordination/communication cost

# Amdahl's Law

- the limit to which we can speed up any complex job is determined by the size of its sequential part
- Speedup  $S_n = \frac{\text{time taken on a single processor}}{\text{time taken on } n \text{ processors}}$
- Amdahl's law characterizes the maximum speedup given  $n$  processors and a parallel fraction  $p$


# Amdahl's Law

$$S_{n,p} = 1 / (1 - p + p / n)$$

Same example: 10 painters, 10 rooms, one room is twice the size of the others:

$$S_{10,10/11} = 1 / (1/11 + 1/11) = 5.5$$

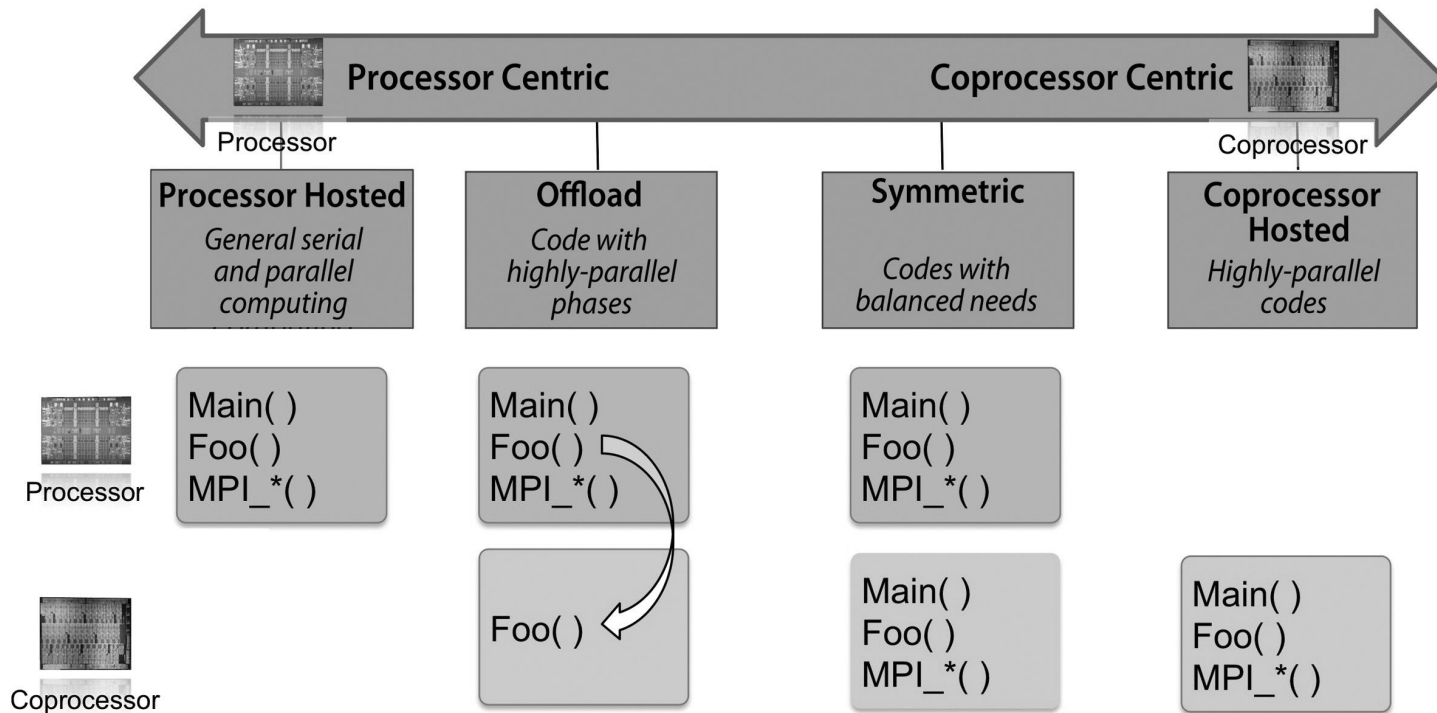
# Amdahls Law



# Conclusion

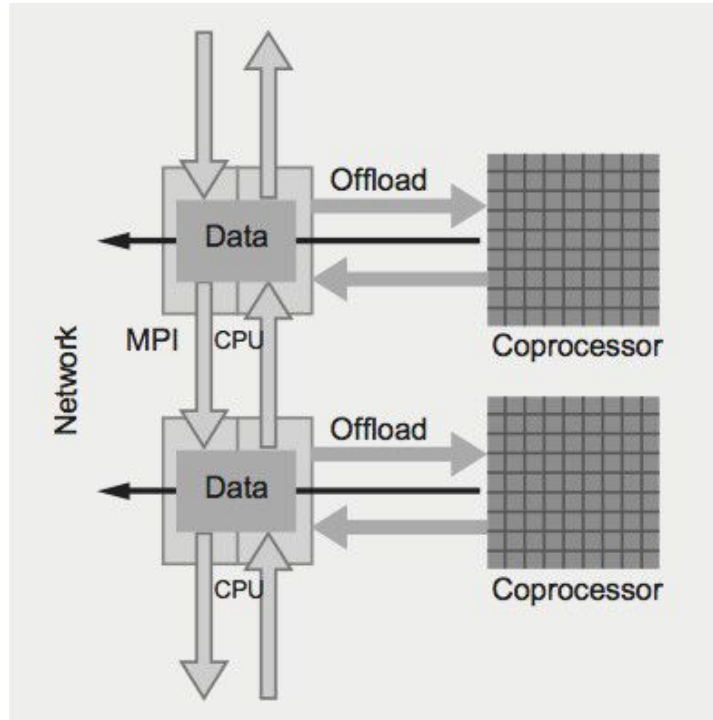
- there will always be a **serial part** of our application
- with the right **tools and engineering principles** we can minimize this serial part
- the eventual **partition** between serial and parallel regions as well as the **hardware setup** at hand will guide the **programming models** we choose

# Programming Model

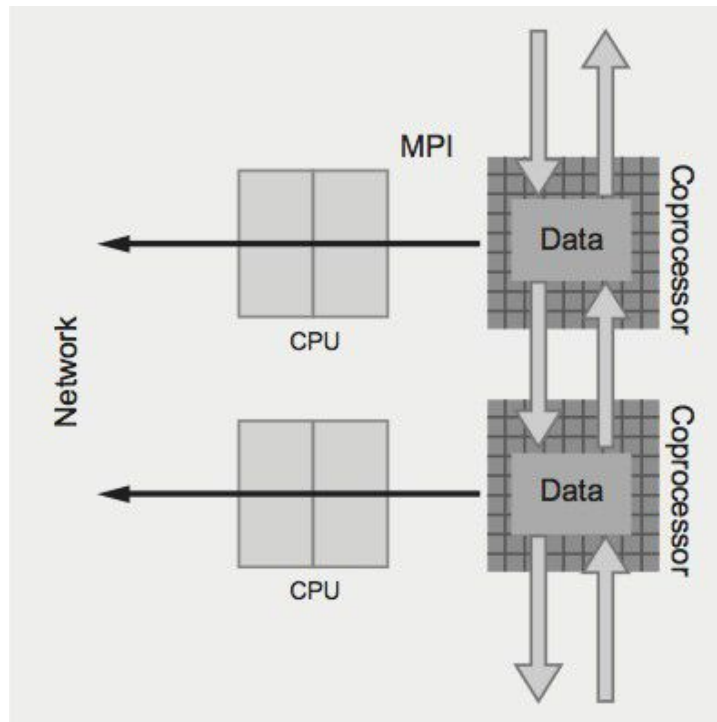




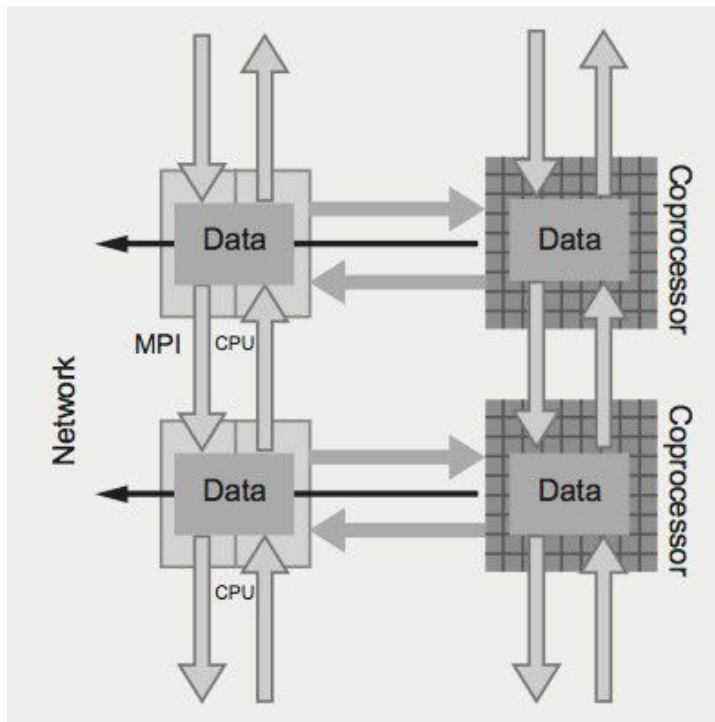
# Offload



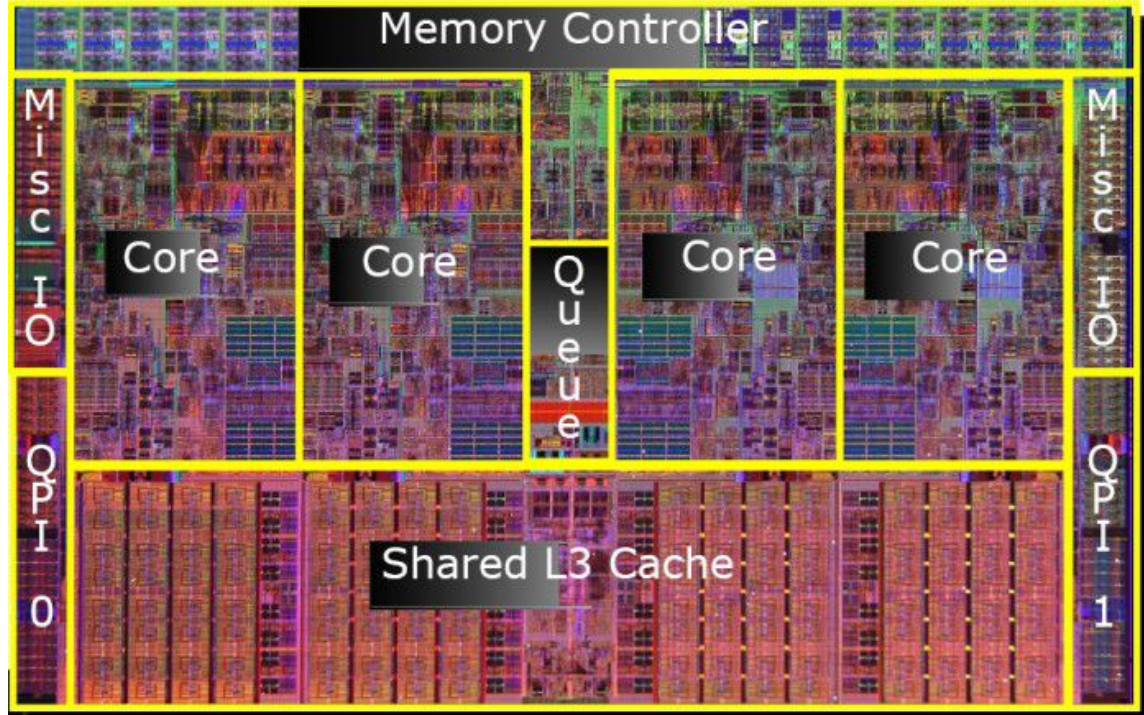
# Native




# Symmetric



# Caches



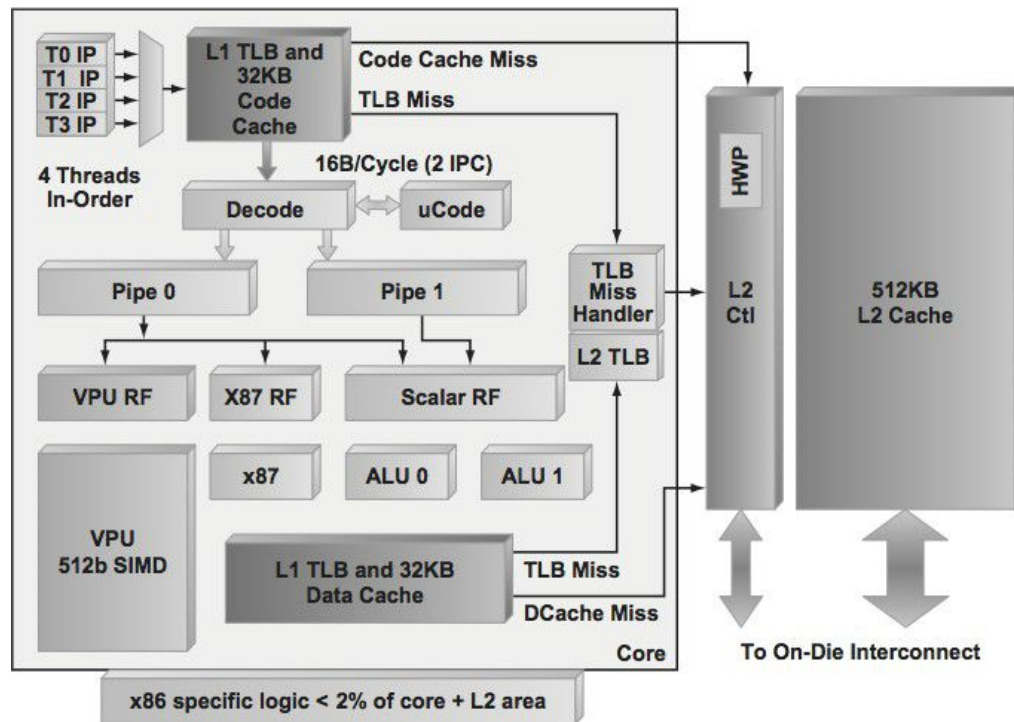
# Caches



# Cache Structure

- modern CPUs are much faster than main memory
  - throughput: 8000 GB/s compute  
~140 GB/s fetch
  - latency: 1 ns compute  
~100 ns cache miss
- orders of magnitude difference
- this resulted in multi-megabyte sized caches

# Cache attached to Core



# Cache Lines

- Data flows among the CPU's caches and memory in fixed-length blocks called **cache lines**
- usually a power of 2, ranging from 16 to 256 bytes
- both the Xeon and Xeon Phi CPU Cores have a cache line size of 64 bytes