


Alignment - Multi-dimensional data structures

```
constexpr int N = 100;
float A[N * N] __attribute__((aligned(64)));
float y[N] __attribute__((aligned(64)));
float x[N] __attribute__((aligned(64)));
for (int i = 0; i < N; ++i) {
    for (int j = 0; j < N; ++j)
        y[i] += A[i * N + j] * x[j];
}
```

Alignment - cont.

- often found in ND-structure (e.g. matrix, $N=2$) traversal
- first element of subsequent rows/columns not aligned to 64 byte boundary!
- this is ignored when `#pragma omp simd` is applied!
→ can lead to crash
- solution: apply **padding** to the row to make it a multiple of 64 bytes in size
- total number of elements allocated for one row/column is called the **leading dimension**

Padding




- well known technique in practice
- extra dummy space to speed-up computation (see false-sharing)
- “wasted”-space can sometimes be used to accommodate for in-loop variables

Strided Access

```
constexpr int N = 100;  
float A[N * N] __attribute__((aligned(64)));  
float y[N] __attribute__((aligned(64)));  
float x[N] __attribute__((aligned(64)));  
for (int j = 0; j < N; ++j) {  
    for (int i = 0; i < N; ++i)  
        y[i] += A[i * N + j] * x[j];  
}
```

Non-unit stride



Strided access

- best performance is achieved when data is accessed with **unit (=1) stride**, i.e. consecutive elements in an array
- a single packed load/store instruction will be generated
- when facing a fixed non-unit stride the compiler might decide to generate **gather/scatter** instructions (slow)
or
non-vectorized code instead (if N is too large)
- **rewrite code** if necessary to avoid non-unit stride

Streaming Stores

- data that the CPU reads/writes needs to be fetched from memory into cache first
- when a memory region is **never read, but only written** to, these elements consume precious cache space
- cache is for **temporal locality**, but there is no reuse in sight!
- streaming stores are a special feature of the Xeon Phi vector units (and upcoming CPUs)
- data that is written and not read is not fetched to cache
- writes go straight to memory over separate buffer
- saves **cache space** and **memory bandwidth**

Example

- z is never read, only written
- add compiler flag `-opt-streaming-stores` [auto|never|always]
- use compiler directive before loop: `#pragma vector nontemporal`

```
constexpr int N = 100;  
float z[N], float x[N], float y[N];  
#pragma vector nontemporal  
for (int i = 0; i < N; ++i)  
    z[i] = x[i] + y[i];
```


Structure of Arrays vs. Array of Structures

- complex data structures usually have multiple attributes combined in a single struct/class → *Object Oriented Programming (OOP) / data encapsulation*
- this can lead to performance penalties due to non-unit stride access
- even if all elements within the structure are processed at the same time (spatial locality), vector units can be left partially empty

SOA vs. AOS

```
constexpr int NUM_PIXELS; = 100;
struct Pixel {
    float r;
    float g;
    float b;
};
// AOS-approach
std::vector<Pixel> pixels(NUM_PIXELS);
// SOA-approach
struct Pixels {
    float r[NUM_PIXELS];
    float g[NUM_PIXELS];
    float b[NUM_PIXELS];
};
```

- if all components of point are visited before going to the next → good locality of reference
- but even then, we either pay for expensive loads due to unaligned structs or the structs are aligned but the vector units are only filled to 75% (at max)
- in case a single member within a struct is visited in a loop, SoA provides better locality of reference AND unit-stride access
- however, it exhibits bad locality of reference when all components of a single struct are required at the same time

SOA vs. AOS - Cont.

Array of Structures

R	G	B	R	G	B	R	G	B
----------	----------	----------	----------	----------	----------	----------	----------	----------

Structure of Arrays

R	R	R	G	G	G	B	B	B
----------	----------	----------	----------	----------	----------	----------	----------	----------

Cache-Blocking

- class of algorithmic techniques
- block data structures to fit in cache by reorganizing memory accesses
- idea: load small subset of larger data set into cache, then work on this **block** while it's in cache
- two types of blocking:
 - using: blocking supports **spatial locality**
 - reusing: blocking supports **temporal locality**
- like all cache-aware techniques: reduces pressure on the main memory bandwidth → speed-up

Blocking-basics

- can be performed both on 1D-data structures or multidimensional data structures
- many applications exhibit multiple accesses to the same memory location
- a cache miss penalty occurs when the data has been evicted from cache when subsequent accesses happen
- blocking usually involves **loop-splitting** and **loop-interchange**
- hard for the compiler to do:
due to strive for correctness → wants to remain loop-iteration order

Blocking-Principle

```
for (int i = 0; i < M; ++i)
  for (int j = 0; j < N; ++j)
    compute(i, j); // some routine using large data
                  // indexed by i and j
```

- for every value of i , compute iterates over the same $0 \leq j < N$ part of the data
- if that part is too big for the machines cache, it will be read from main memory M times
- use blocking, if the order of the $M * N$ iterations doesn't matter

Loop-Splitting (1)

```
constexpr int BLOCK_SIZE = 1;  
for (int ii = 0; ii < M; ii+=BLOCK_SIZE)  
    for (int i = ii; i < ii + BLOCK_SIZE; ++i)  
        for (int jj = 0; jj < N; jj+=BLOCKS_SIZE)  
            for (int j = jj; j < jj + BLOCK_SIZE; ++j)  
                compute(i, j);
```

Loop-Splitting (2)

```
constexpr int BLOCK_SIZE = 128;  
using namespace std;  
for (int ii = 0; ii < M; ii+=BLOCK_SIZE)  
    for (int i = ii; i < min(ii + BLOCK_SIZE, M); ++i)  
        for (int jj = 0; jj < N; jj+=BLOCKS_SIZE)  
            for (int j = jj; j < min(jj + BLOCK_SIZE, N); ++j)  
                compute(i, j);
```


Loop-Interchange


```
constexpr int BLOCK_SIZE = 128;  
using namespace std;  
for (int ii = 0; ii < M; ii+=BLOCK_SIZE)  
    for (int jj = 0; jj < N; jj+=BLOCKS_SIZE)  
        for (int i = ii; i < min(ii + BLOCK_SIZE, M); ++i)  
            for (int j = jj; j < min(jj + BLOCK_SIZE, N); ++j)  
                compute(i, j);
```

- now the access pattern changed
- i and j are now confined to the respective ii-jj-block defined by BLOCK_SIZE

About BLOCK_SIZE

- the BLOCK_SIZE parameter **needs to be tuned** for the particular application and target machine
- there can be **multiple levels** of blocking for multiple memory hierarchy transitions, e.g. blocking on L2 cache and blocking on L1 cache
- the BLOCK_SIZE does not have to be the same for every dimension → blocks can be **rectangular**: this might help for better cache utilization

Blocking vs Non-Blocking




Example: Matrix Transposition

// assume square matrix for simplicity

```
void transpose(double * m, const int N) {  
    for (int i = 0; i < N; ++i)  
        for (int j = 0; j < N; ++j)  
            std::swap(m[i * N + j], m[j * N + i]);  
}
```

- every element is accessed at most once → no temporal reuse possible
- big rows/columns might not fit into cache and evict cache lines containing next element
- improve **spatial locality**

Example: Matrix Multiplication



- conceptually, B is looped over m times
- If B is too big for the machine's cache, B will have to be read m times from main memory
- Blocking can improve [temporal locality](#)

Vectorization guidelines

- develop a habit to think about your access pattern even **before** you start to vectorize
- make your data-structures vector-friendly:
usually static/dynamic **arrays**
- make sure your inner-most loop-index corresponds to **unit stride access** index of array
- prefer structure of array (SoA) over array of structures (AoS), if(!) your access pattern would benefit from **vectorization**