

Synchronization

- multiple threads reading a shared memory location is not problematic
- when writing to a single memory item we need synchronization to avoid **data races**
- OpenMP provides mainly two means to isolate read/write accesses to variables **critical** and **atomic** directives



Example

```
int N = 1000000000;  
std::vector<int> prime_numbers;  
#pragma omp parallel for schedule(dynamic)  
for (int i = 0; i < N; ++i) {  
    if (is_prime(i))  
        prime_numbers.push_back(i);  
}
```

- multiple threads alter internal state of std::vector concurrently
- the STL in general, and std::vector in particular are not thread-safe
- we need to make sure **only one thread at a time** adds an element to the vector



Synchronization: critical-directive

- a region of code that must be executed by only one thread at a time
- `#pragma omp critical [(name)]`
- critical sections with the same name are treated as the same protected section
- when no name is given, critical sections belong to the **global** name
- give your critical sections meaningful names according to their semantic



Example

```
int N = 10000000;
std::vector<int> prime_numbers;
#pragma omp parallel for schedule(dynamic)
for (int i = 0; i < N; ++i)
    if (is_prime(i))
        #pragma omp critical (prime_insert)
        prime_numbers.push_back(i);
```



What does critical do internally?

```
std::mutex prime_insert;
int N = 1000000000;
std::vector<int> prime_numbers;
#pragma omp parallel for schedule(dynamic)
for (int i = 0; i < N; ++i) {
    if (is_prime(i)) {
        std::lock_guard<std::mutex> prime_insert_lock(prime_insert);
        prime_numbers.push_back(i);
    }
}
```

- all major OpenMP implementations use locks to implement the critical-directive
- locks are expensive: **use wisely and sparingly**
- can destroy performance in tight loops on frequently updated data-structures



Synchronization: atomic-directive

- often a cheaper alternative to expensive locks
- `#pragma omp atomic`
- can be applied only to certain **binary operations**
+, -, *, /, shift- and logic operators
- mostly used to increment/decrement a variable
- may not be available on all architectures
- requires **special hardware support**
- basic building block for **lock-free programming**



Example

```
int N = 10000000000;  
int num_primes = 0;  
#pragma omp parallel for schedule(dynamic)  
for (int i = 0; i < N; ++i) {  
    if (is_prime(i))  
        #pragma omp atomic  
        ++num_primes;  
}
```

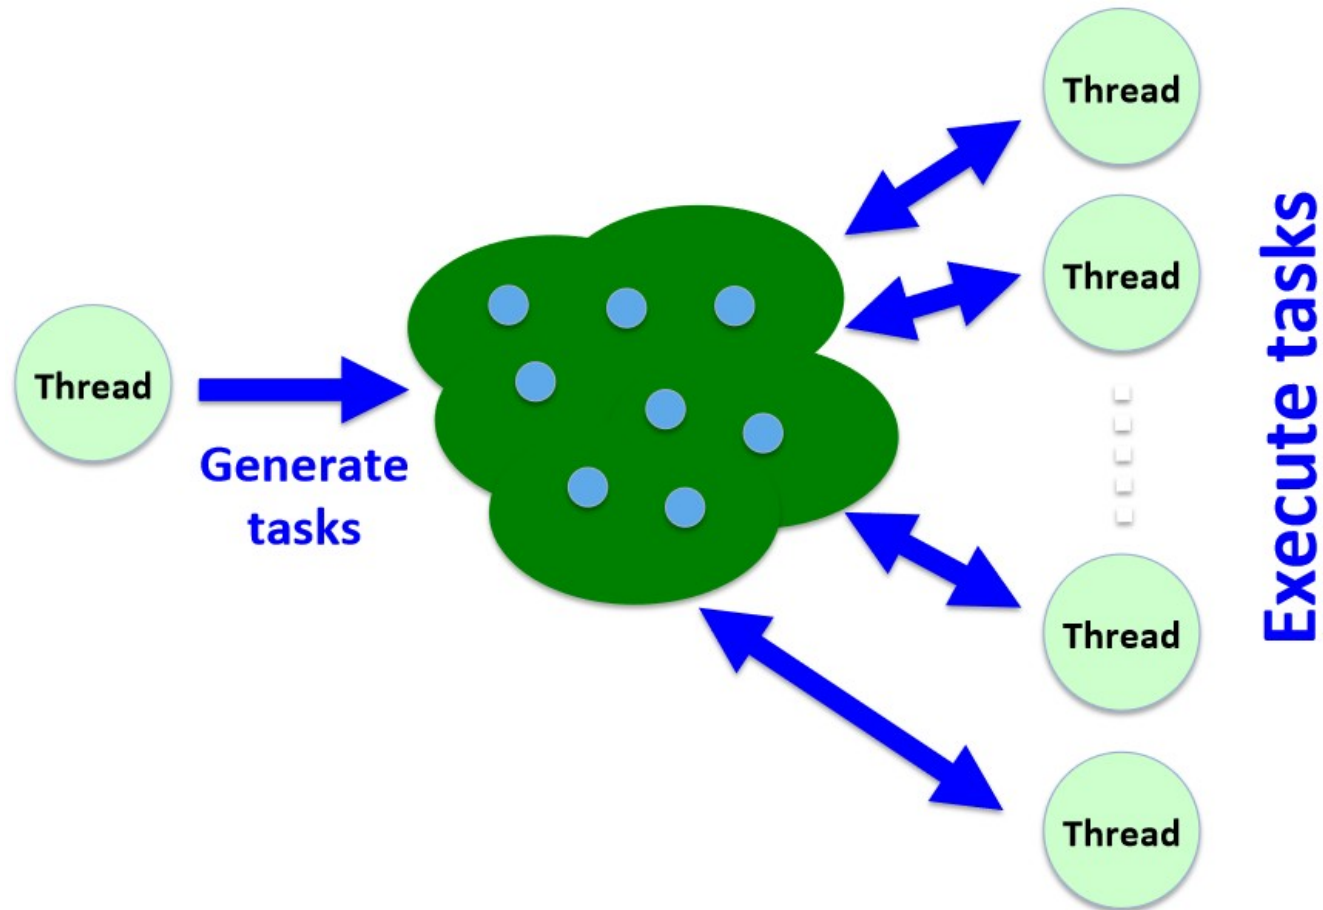


OpenMP Tasking

- up until OpenMP 2.5 directives were directed towards more **regular** program structure
 - loop iterations known at runtime (**fixed!**)
 - only a **finite** number of parallel regions (nesting + finite thread number)
- not suitable for **linked lists** or **recursive** algorithms (possible, but ugly at best)
- OpenMP 3.0 Tasks: good for **irregular** problems
- **tasks** are lightweight encapsulations of work



Tasking concept



OpenMP task directive

- `#pragma omp task [clauses]`
- known clauses: **default**, **private**, **firstprivate**, **shared**, **if**
- new: **final**, **untied**, **mergeable**
- task is bound to the innermost enclosing parallel region and its thread team
- the tasks code is the following structured block together with a **data-environment** according to the usual data sharing rules
- tasks can be executed immediately or **deferred**



task clauses

- **if(expr)** : if expression evaluates to false, no task is generated, the current thread executed the code immediately -> performance optimization
- **final(expr)**: similar to if, but all child tasks inherit the final property (e.g. recursive algorithms reached certain depth)
- **untied**: if executing thread is suspended (for whatever reason), another thread may “steal” the task and continue
- **mergeable**: not really beneficial in practice: allow the runtime to merge the tasks data environment with its calling environment



Example

```
#pragma omp parallel
{
    #pragma omp single
    while (my_pointer) {
        #pragma omp task firstprivate(my_pointer)
        (void) do_independent_work(my_pointer); // the task's code
        my_pointer = my_pointer->next;
    } // implicit end of single
} // end of parallel region
```



When do task get executed?

- depending on state: immediately or deferred
- **immediate**
if-clause evaluate to false or final-clause evaluate to true
- deferred tasks are executed at **barriers** or **taskwait** constructs
 - `#pragma omp barrier` and all implicit barriers (end of parallel region, for loop sharing, ...)
 - `#pragma omp taskwait`



taskyield directive

- performance hint to the OpenMP runtime
- suspend the current task to allow the executing thread to do other (useful) work

```
#include <omp.h>
void something_useful();
void something_critical();
void foo(omp_lock_t * lock, int n) {
    for(int i = 0; i < n; i++)
        #pragma omp task
        {
            something_useful();
            while( !omp_test_lock(lock) ) {
                #pragma omp taskyield
            }
            something_critical();
            omp_unset_lock(lock);
        }
}
```



OpenMP environment

- Variables that control the runtime behavior of an application

- **OMP_NUM_THREADS**

sets the maximum default number of threads in a team

```
OMP_NUM_THREADS = 2 ./program
```

no recompilation needed

- **OMP_PROC_BIND**

tell the runtime to pin threads to specific cores

no context switch and hot caches

```
OMP_PROC_BIND=true ./program
```

- **OMP_NESTED**

activates nested parallelism, which is deactivated by default

```
OMP_NESTED = true ./program
```

